

CMAD Adv I

Polyglot Databases

- Understand Consistency Models - ACID/BASE
- Whats driving us to this
- Benefits
- Challenges
- Deciding Criteria

ACID/BASE

- Limits of instantaneous consistency
 - The balance between write performance and availability in instantaneous consistent systems
- What does Eventual Consistency get us
- CAP Theorem

Deciding Criteria

- Write performance
- Read performance
- Availability
- Consistency
- Partition Tolerance
- Query patterns
 - Text search, graph search

Types of DBs

- Mongo
 - Tuneable consistency, Availability, Partition Tolerance.
 - Indexed Document Store
 - Master slave model with writes only to master - Easy scalability for reads
- Neo4j
 - Graph database tuned for graph questions
 - Consistent and Available. Billions of nodes and relationships

Types of DB

- Cassandra
 - Column store with multi master ability. Closest to SQL
 - Available and Partition Tolerant
 - Easy scalability for writes
- Redis in-memory key-value store

Question of Durability

- What happens to backups
- How facebook handles redundancy and durability

Threading and Locking

- Understand thread racing
- Solving thread racing with locks
- Using atomic variables

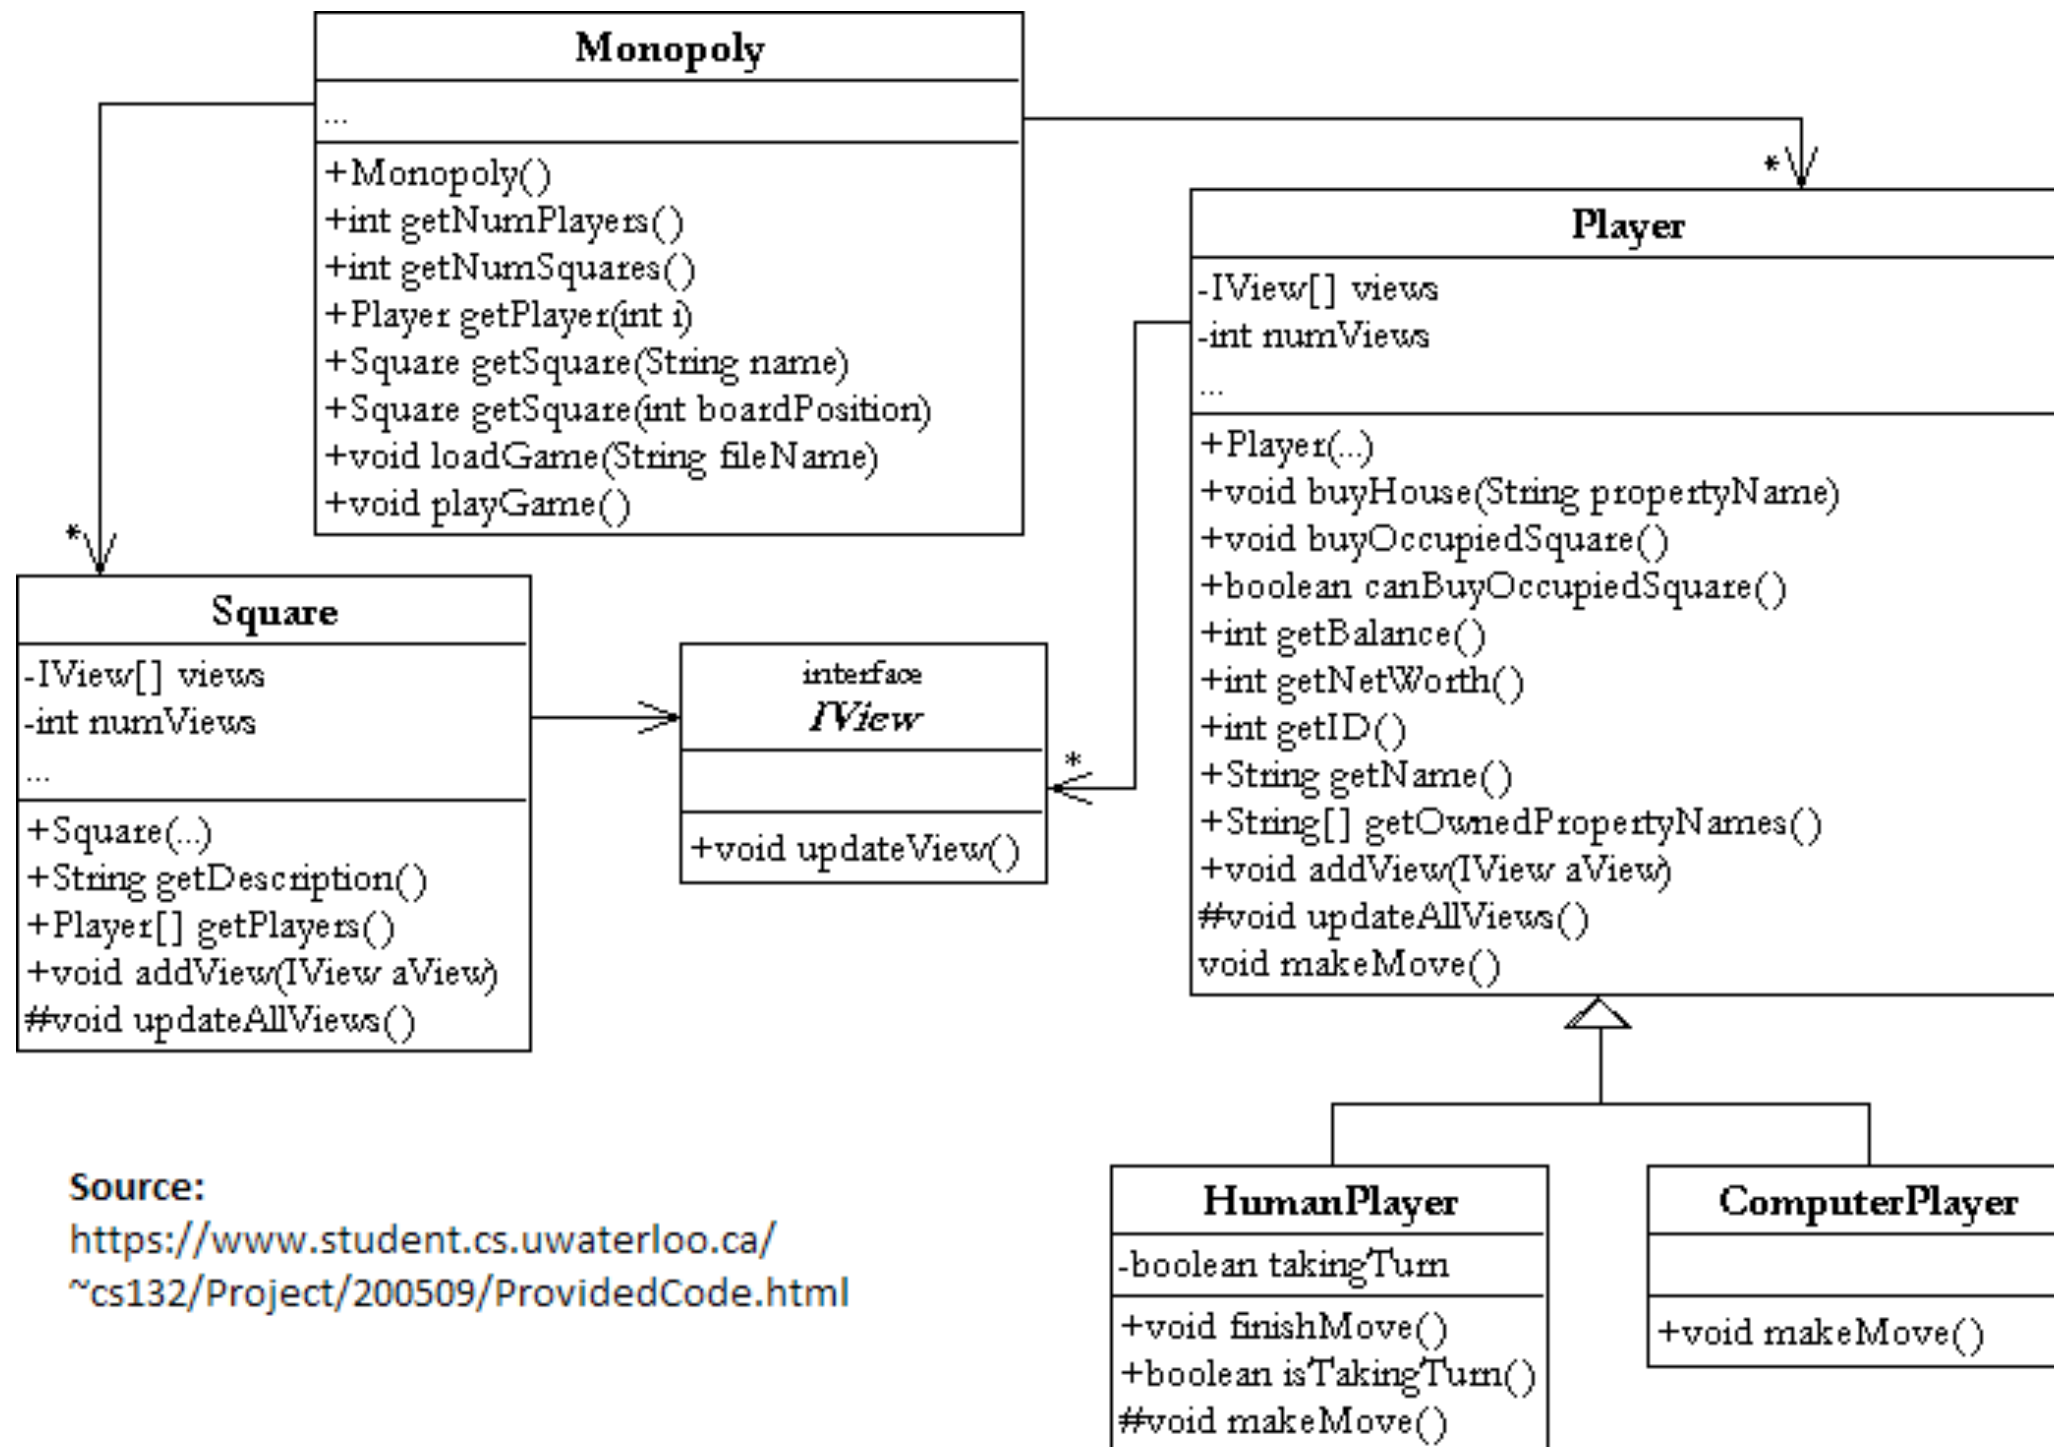
Shared Collections

- Using shared collections
 - Write and read operations. Try multiple thread read and write into a synchronized map
 - Compare it with concurrent map implementation
- Fail fast and snapshot iterators
- CopyOnWriteArrayLists

I/O Performance

- Java IO library read operation
- Asynchronous Read using NIO
- Measure performance in multi threaded environment

Traditional OO Design



Source:

<https://www.student.cs.uwaterloo.ca/~cs132/Project/200509/ProvidedCode.html>

Stateless OO Design Classes

- Model Classes - No patterns, just DTOs
- Business Classes - Design patterns apply here
- Technology Classes
- Utility Classes

Vert.x


- We can build reactive, non blocking, event driven apps
- Polyglot - javascript, java, groovy and ruby
- Not an app server, modular
- Good fit for microservices

Vert.x Performance

Best plaintext responses per second, i7-2600K hardware (55 tests)										
Framework	Best performance (higher is better)		Cls	Lng	Plt	FE	Aos	IA	Errors	
vertx	656,119	<div><div></div></div> 100.0%	Plt	Jav	vtx	Non	Lin	Rea	0	
netty	632,596	<div><div></div></div> 96.4%	Plt	Jav	Nty	Non	Lin	Rea	10	
undertow	614,547	<div><div></div></div> 93.7%	Plt	Jav	Utw	Non	Lin	Rea	70	
plain-windows	611,095	<div><div></div></div> 93.1%	Ful	Sca	Pla	Non	Win	Rea	52	
plain	543,670	<div><div></div></div> 82.9%	Ful	Sca	Pla	Non	Lin	Rea	21,473	
cpoll_cppsp	522,423	<div><div></div></div> 79.6%	Mcr	C++	Cpl	Non	Lin	Rea	0	
plain-servlet-linux	450,462	<div><div></div></div> 68.7%	Ful	Sca	Pla	Non	Lin	Rea	3,406	
jetty-servlet	448,947	<div><div></div></div> 68.4%	Plt	Jav	Jty	Jty	Lin	Rea	351	
grizzly	425,172	<div><div></div></div> 64.8%	Mcr	Jav	Svt	Grz	Lin	Rea	149	
gemini	424,586	<div><div></div></div> 64.7%	Ful	Jav	Svt	Non	Lin	Rea	2,071	
spray	422,431	<div><div></div></div> 64.4%	Mcr	Sca	Spr	Non	Lin	Rea	1,003	
servlet	417,619	<div><div></div></div> 63.6%	Plt	Jav	Svt	Res	Lin	Rea	1,214	
go	367,274	<div><div></div></div> 56.0%	Plt	Go	Go	Non	Lin	Rea	230	
openresty	330,829	<div><div></div></div> 50.4%	Plt	Lua	OpR	ngx	Lin	Rea	1,517	
spark	226,365	<div><div></div></div> 34.5%	Mcr	Jav	Svt	Res	Lin	Rea	0	
revel	221,603	<div><div></div></div> 33.8%	Ful	Go	Go	Non	Lin	Rea	521	
falcore	172,827	<div><div></div></div> 26.3%	Mcr	Go	Go	Non	Lin	Rea	0	
compojure	127,671	<div><div></div></div> 19.5%	Mcr	Clj	Svt	Res	Lin	Rea	0	
http-listener	123,434	<div><div></div></div> 18.8%	Plt	C#	Net	hts	Win	Rea	0	
falcon-pypy	112,550	<div><div></div></div> 17.2%	Mcr	Py	Gun	Non	Lin	Rea	0	
wsgi-nginx-uwsgi	110,522	<div><div></div></div> 16.8%	Plt	Py	uWS	ngx	Lin	Rea	459,212	
evhttp-sharp	107,494	<div><div></div></div> 16.4%	Mcr	C#	Mon	Non	Lin	Rea	0	
bottle-pypy	88,328	<div><div></div></div> 13.5%	Mcr	Py	Tor	Non	Lin	Rea	52	
nodejs	80,363	<div><div></div></div> 12.2%	Plt	JS	njs	Non	Lin	Rea	0	

Verticles

- Vert.x can deploy and execute components called Verticles.
- You can think of verticles as being similar to servlets or message driven EJBs driven by an event bus

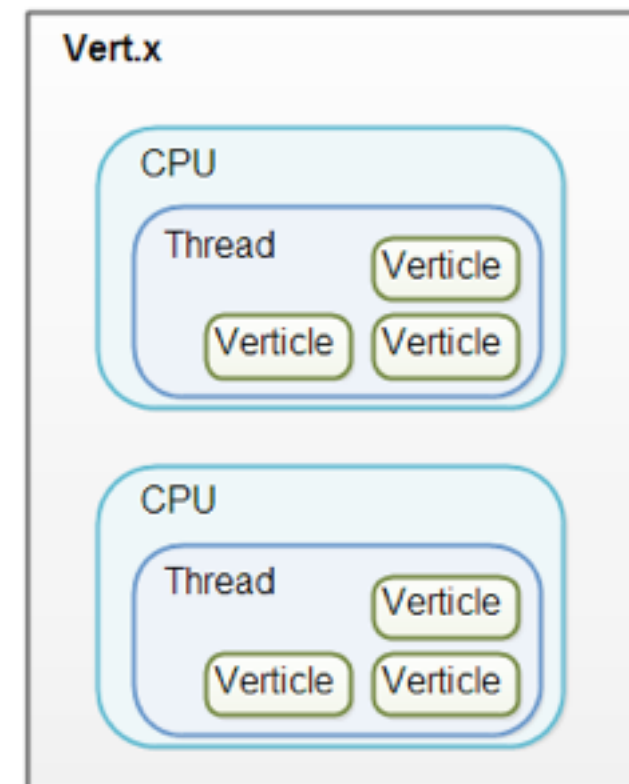


Verticle Messaging

- Messages can be simple objects (e.g. Java objects), strings, CSV, JSON, binary data or whatever else you need
- Verticles can send and listen to addresses. An address is like a named channel.
- When a message is sent to a given address, all verticles that listen on that address receive the message.
- Verticles can subscribe and unsubscribe to addresses without the senders knowing.

Threading Model

- verticle is only ever executed by a single thread, and always by the same thread.
- A single thread can distribute messages to multiple verticles.
- Vert.x creates one thread per CPU
- Vert.x comes with a set of built-in services (functionality).
Some of these services are:
 - HTTP server
 - JDBC connector
 - MongoDB connector



Number of Instances

- Deploy multiple verticle instances to run in different threads
- `DeploymentOptions options = new DeploymentOptions().setInstances(16);`

Creating Verticles

```
public class MyVerticle extends AbstractVerticle {

    @Override
    public void start(Future<Void> startFuture) {
        System.out.println("MyVerticle started!");
    }

    @Override
    public void stop(Future stopFuture) throws Exception {
        System.out.println("MyVerticle stopped!");
    }

}

public static void main(String[] args) {
    VertxOptions options = new VertxOptions().setWorkerPoolSize(10);
    Vertx vertx = Vertx.vertx(options);
    vertx.deployVerticle("com.mydomain.MyVerticle");
}
```

Verticle Events

- The start() method : start HTTP or TCP server, register event handlers on the event bus, deploy other verticles, or whatever else your verticle needs to do its work.

- Shutdown stuff in stop method

- The verticle will be deployed asynchronously

```
vertx.deployVerticle("com.mydomain.MyVerticle", new  
Handler<AsyncResult<String>>() {  
    @Override  
    public void handle(AsyncResult<String> stringAsyncResult) {  
        System.out.println("Verticle deployment complete");  
    }  
});
```

Registering to Events

- When a verticle wants to listen for messages from the event bus, it listens on a certain address. An address is just a name (a String) which you can choose freely.
- An address is thus more like the name of a channel with multiple receivers

```
vertx.eventBus().consumer("Channel1", message -> {  
    System.out.println("message.body() = "  
        + message.body());  
});  
vertx.eventBus().publish("Channel1", "message 2");
```

Types Of Verticles

- Standard Verticles
 - These are the most common and useful type - they are always executed using an event loop thread.
- Worker Verticles
 - These run using a thread from the worker pool. An instance is never executed concurrently by more than one thread.

```
DeploymentOptions options = new DeploymentOptions().setWorker(true);  
vertx.deployVerticle("com.mycompany.MyOrderProcessorVerticle", options);
```

- Multi-threaded worker verticles
 - These run using a thread from the worker pool. An instance can be executed concurrently by more than one thread.

Vert.x buffers

- Carry Binary Information in buffers. Dynamically resizable

- Can be used as message payloads

```
byte[] initialData = new byte[]{1, 2, 3};
```

```
Buffer buffer = Buffer.buffer(initialData);
```

```
buffer.setShort ( 10, (short) 127);
```

```
buffer.appendByte ((byte) 127);
```

Running Blocking Code

- When we HAVE to invoke synchronous APIs, vertx provides a way to do that:

```
vertx.executeBlocking(future -> {  
    // Call some blocking API that takes a significant amount of time  
    // to return  
    String result = someAPI.blockingMethod("hello");  
    future.complete(result);  
}, res -> {  
    System.out.println("The result is: " + res.result());  
});
```


Mongodb

- Mongodb is an indexed document store.
- A document is typically something like a json structure
- Can be indexed based on any of the fields
- Can be replicated, sharded, tunable consistency with single master and leader election.

- `db.users.insert({name:'Hari',email:'hari@abc.com'})`
- `show collections`
- `db.users.update({age: { $gt: 18 }},{ $set: {status:'A'}},{multi: true})`
- `db.users.remove({status: 'D'})`

Mongodb Queries

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

```
SELECT _id, name, address  
FROM users  
WHERE age > 18  
LIMIT 5
```

← projection
← table
← select criteria
← cursor modifier

Creating Index

- `db.users.createIndex({ email: 1 })`
- `db.users.createIndex({ email: 1, name: 1 })`
- Fully covered queries:
 - `db.users.find { email: /. *yahoo.com/},{ name: 1, _id: 0 })`


Integrating With Java

- POM dependencies

Neo4j

- ACID transactions,
- High availability,
- Scales to billions of nodes and relationships,
- High speed querying through traversals,
- Declarative graph query language.

Graph



Features

- SQL Like easy query language Neo4j CQL
- It supports Indexes by using Apache Lucence
- It supports UNIQUE constraints
- It contains a UI to execute CQL Commands : Neo4j Data Browser
- It supports full ACID
- It uses Native graph storage with Native GPE(Graph Processing Engine)
- It supports exporting of query data to JSON and XLS format

Features

- It provides REST API for the data
- It supports two kinds of Java API: Cypher API and Native Java API to develop Java applications.

CQL

- `create (blog:Blog)` - node name blog, label Blog
- `create (:Blog{title:'Some blog',desc:'This is some blog'})` - Create with properties
- `match(b:Blog) return b` - select * from blog b
- `MATCH (b:Blog),(u:User) CREATE (b)-[r:AUTHORED_BY]->(u)`
 - Form relationship between EVERY blog and EVERY User

CQL

- MATCH (u:User{name:'Faizal'}) CREATE (blog:Blog{title:'related blog',desc:'This is some related blog'})-[r:AUTHORED_BY{when:'5th Sept'}]->(u)
 - Create an object with a relationship
- match(b:Blog)-[a:AUTHORED_BY]-(u:User{name:'Faizal'}) return a.when
 - Find all dates when Faizal authored blogs
- MATCH (b:Blog) WHERE b.title =~ '.*Some.*' RETURN b

CQL

- MATCH (b:Blog)-[rel]-(u:User) DELETE rel
 - Delete only the relationships
- MATCH (u:User{name:'Faizal'}) SET u.sex = 'Male'
RETURN u
 - Set a field and return the node

Java API support

- Neo4J supports two types of Java APIs
 - Native Java API
 - CQL API